

# Linear Regression

## Univariate linear regression and gradient descent

Khiem Nguyen

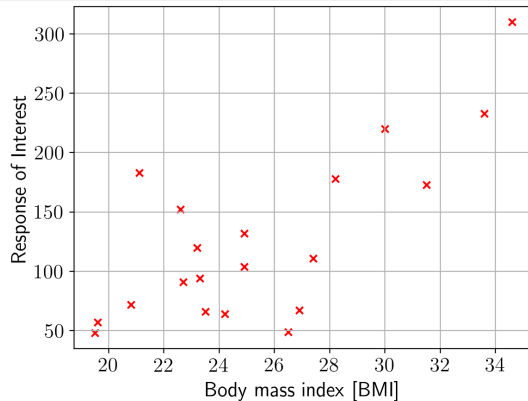
Email	<code>khiem.nguyen@glasgow.ac.uk</code>
MS Teams	<code>khiem.nguyen@glasgow.ac.uk</code>
Whatsapp	+44 7729 532071 (Emergency only)

May 18, 2025

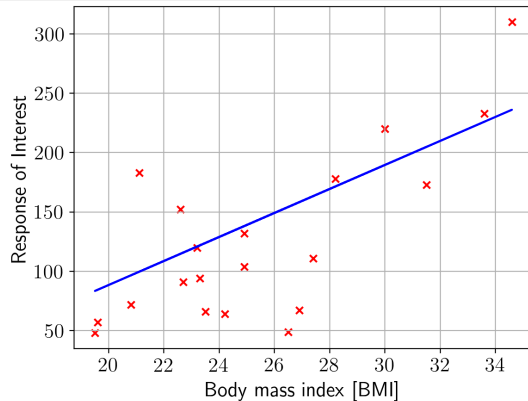


University  
of Glasgow


# Linear regression: presentation



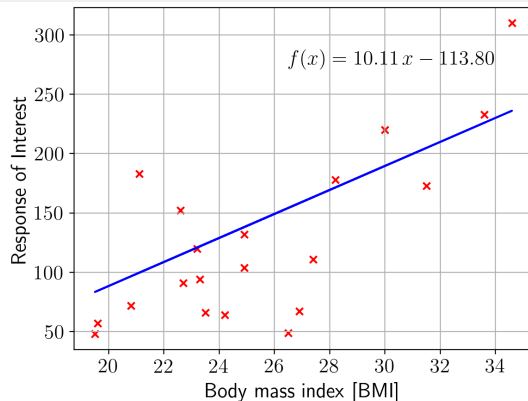
# Linear regression: presentation



# Linear regression: presentation




# Linear regression: presentation



- Supervised learning model data has “right answers”.
- Regression model predicts numbers.
- Classification model predicts categories.

## Linear regression: Data table



BMI	RI
32.1	151
21.6	75
30.5	141
25.3	206
...	...
23	135

# Terminology

Training	Data used to train the model	
	$x$	$y$
	BMI	RI
(1)	32.1	151
(2)	21.6	75
(3)	30.5	141
(4)	25.3	206
...	...	...
(m)	23	135

# Terminology

Training	Data used to train the model	
	$x$	$y$
	BMI	RI
(1)	32.1	151
(2)	21.6	75
(3)	30.5	141
(4)	25.3	206
...	...	...
(m)	23	135

## Notation

$x$  = input variable

feature

$y$  = output variable

target variable



# Terminology

Training	Data used to train the model	
	$x$	$y$
	BMI	RI
(1)	32.1	151
(2)	21.6	75
(3)	30.5	141
(4)	25.3	206
...	...	...
(m)	23	135

## Notation

$x$  = input variable

feature

$y$  = output variable

target variable

$(x, y)$  = single training example

$(x^{(i)}, y^{(i)})$  =  $i^{\text{th}}$  training example

We have  $(1^{\text{st}}, 2^{\text{nd}}, 3^{\text{rd}}, \dots, N^{\text{th}})$

# Terminology

Training	Data used to train the model	
	$x$	$y$
	BMI	RI
(1)	32.1	151
(2)	21.6	75
(3)	30.5	141
(4)	25.3	206
...	...	...
(m)	23	135

$$x^{(1)} = 32.1, \quad y^{(1)} = 151$$

$$x^{(2)} = 21.6, \quad y^{(2)} = 75$$

$$(x^{(1)}, y^{(1)}) = (32.1, 151)$$

Note:  $x^{(2)} \neq x^2 \rightarrow$  not exponent, just indexing

## Notation

$x$  = input variable

feature

$y$  = output variable

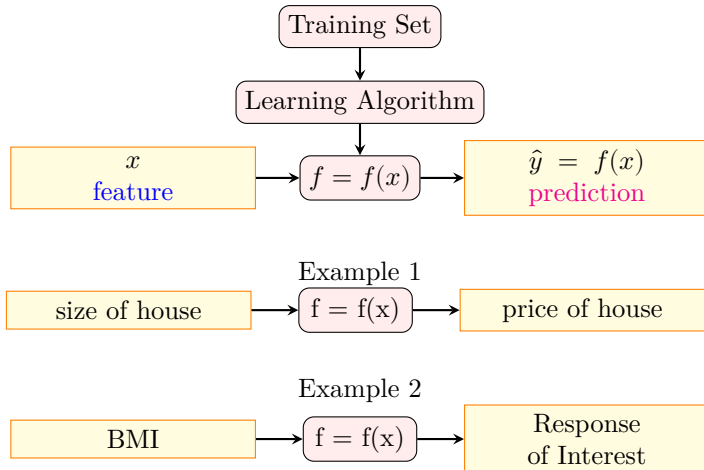
target variable

$(x, y)$  = single training example

$(x^{(i)}, y^{(i)}) = i^{\text{th}}$  training example


We have  $(1^{\text{st}}, 2^{\text{nd}}, 3^{\text{rd}}, \dots, N^{\text{th}})$

# Learning a hypothesis/function model




**Linear function:**  $f(x) = wx + b, \quad w, b \in \mathbb{R}$

# Learning a hypothesis/function model



$$f(x) = f_{w,b}(x) = wx + b \quad (1)$$



- Linear regression with one variable (single feature  $x$ )
- Univariate linear regression  
one variable

## Interpretation of linear model

$x$	$y$
32.1	151
21.6	75
30.5	141
25.3	206
...	...
23	135

# Interpretation of linear model

Model

$$f_{w,b}(x) = wx + b, \quad w, b \text{ -- parameters}$$

$x$	$y$
32.1	151
21.6	75
30.5	141
25.3	206
...	...
23	135

# Interpretation of linear model

Model


$$f_{w,b}(x) = wx + b, \quad w, b \text{ -- parameters}$$

What do the beasts  $w, b$  do, and  
Where (how) to find them?

$x$	$y$
32.1	151
21.6	75
30.5	141
25.3	206
...	...
23	135




## Cost function: detailed explanation







## Cost function: detailed explanation



## Cost function: detailed explanation




## Cost function: detailed explanation




Minimize the total distance  $[d_1(w, b) + d_2(w, b)]$  with respect to  $w, b$

$$\min_{w,b} [d_1(w, b) + d_2(w, b)]$$

## Cost function: detailed explanation




## Cost function: detailed explanation



Many more data points  $\rightarrow$  many more distances  $d_j = d_j(w, b)$ ,  $j = 1, \dots, m$


## Cost function: detailed explanation



Many more data points  $\rightarrow$  many more distances  $d_j = d_j(w, b)$ ,  $j = 1, \dots, m$

$$\min_{w, b} \left\{ \left[ d_1(w, b) + \dots + d_m(w, b) \right] \right\}$$


## Cost function: detailed explanation



Many more data points  $\rightarrow$  many more distances  $d_j = d_j(w, b)$ ,  $j = 1, \dots, m$

$$\min_{w, b} \left\{ \frac{1}{m} \left[ d_1(w, b) + \dots + d_m(w, b) \right] \right\}$$

## Cost function: detailed explanation




Many more data points  $\rightarrow$  many more distances  $d_j = d_j(w, b)$ ,  $j = 1, \dots, m$

$$\min_{w, b} \left\{ \frac{1}{m} \left[ d_1(w, b) + \dots + d_m(w, b) \right] \right\}$$



## Cost function: detailed explanation



Many more data points  $\rightarrow$  many more distances  $d_j = d_j(w, b)$ ,  $j = 1, \dots, m$

$$\min_{w, b} \left\{ \frac{1}{m} \sum_{i=1}^m d_i(w, b) \right\}$$

## Cost function: detailed explanation

But ... life is always more complicated than it looks ☺

## Cost function: detailed explanation

But ... life is always more complicated than it looks ☺

Let us look at three data points  $\{\mathbf{p}^{(1)} = (1, 3), \mathbf{p}^{(2)} = (2, 4), \mathbf{p}^{(3)} = (3, 5)\}$ :

$$\min_{w,b} G(w,b) = \min_{w,b} \frac{1}{3} (|w \times 1 + b - 3| + |w \times 2 + b - 4| + |w \times 3 - 5|)$$

## Cost function: detailed explanation

But ... life is always more complicated than it looks ☺

Let us look at three data points  $\{\mathbf{p}^{(1)} = (1, 3), \mathbf{p}^{(2)} = (2, 4), \mathbf{p}^{(3)} = (3, 5)\}$ :

$$\min_{w,b} G(w,b) = \min_{w,b} \frac{1}{3} (|w \times 1 + b - 3| + |w \times 2 + b - 4| + |w \times 3 - 5|)$$

- $G(w,b)$  is just linear in both  $w$  and  $b$  depending on the different regions
- We have to consider different regions to remove absolutes

## Cost function: detailed explanation

But ... life is always more complicated than it looks ☺

Let us look at three data points  $\{\mathbf{p}^{(1)} = (1, 3), \mathbf{p}^{(2)} = (2, 4), \mathbf{p}^{(3)} = (3, 5)\}$ :

$$\min_{w,b} G(w,b) = \min_{w,b} \frac{1}{3} (|w \times 1 + b - 3| + |w \times 2 + b - 4| + |w \times 3 - 5|)$$

- $G(w,b)$  is just linear in both  $w$  and  $b$  depending on the different regions
- We have to consider different regions to remove absolutes

But ... but how about this?

$$\min_{w,b} \mathcal{L}(w,b) = \frac{1}{3} [(w \cdot 1 + b - 3)^2 + (w \cdot 2 + b - 4)^2 + (w \cdot 3 + b - 5)^2] \quad (2)$$

- $\mathcal{L}(w,b)$  is a quadratic function in both  $w$  and  $b$
- Finding the minimum of quadratic function is easy.

## Cost function: Squared error cost function

$$\begin{aligned}\hat{y}^{(i)} &= f_{w,b}(x^{(i)}) \\ &= wx^{(i)} + b\end{aligned}$$

$$\mathcal{L}(w, b) =$$

=

## Cost function: Squared error cost function

$$\begin{aligned}\hat{y}^{(i)} &= f_{w,b}(x^{(i)}) \\ &= wx^{(i)} + b\end{aligned}$$

$$\begin{aligned}\mathcal{L}(w, b) &= \left(\hat{y}^{(i)} - y^{(i)}\right)^2 \\ &= \end{aligned}$$

## Cost function: Squared error cost function

$$\begin{aligned}\hat{y}^{(i)} &= f_{w,b}(x^{(i)}) \\ &= wx^{(i)} + b\end{aligned}$$

$$\begin{aligned}\mathcal{L}(w, b) &= \frac{1}{2m} \sum_{i=1}^m \left( \hat{y}^{(i)} - y^{(i)} \right)^2 \\ &= \end{aligned}$$



## Cost function: Squared error cost function

$$\begin{aligned}\hat{y}^{(i)} &= f_{w,b}(x^{(i)}) \\ &= wx^{(i)} + b\end{aligned}$$

$$\begin{aligned}\mathcal{L}(w, b) &= \frac{1}{2m} \sum_{i=1}^m \left( \hat{y}^{(i)} - y^{(i)} \right)^2 \\ &= \frac{1}{2m} \sum_i \left( f_{w,b}(x^{(i)}) - y^{(i)} \right)^2\end{aligned}$$

## Cost function: Squared error cost function

$$\begin{aligned}\hat{y}^{(i)} &= f_{w,b}(x^{(i)}) \\ &= wx^{(i)} + b\end{aligned}$$

$$\begin{aligned}\mathcal{L}(w, b) &= \frac{1}{2m} \sum_{i=1}^m \left( \hat{y}^{(i)} - y^{(i)} \right)^2 \\ &= \frac{1}{2m} \sum_i \left( f_{w,b}(x^{(i)}) - y^{(i)} \right)^2\end{aligned}$$

Find  $w, b$  so that:  $\hat{y}^{(i)}$  is close to  $y^{(i)}$  for all data points  $(x^{(i)}, y^{(i)})$ .

$$\min_{w,b} \mathcal{L}(w, b)$$

## Cost function: Intuition (again!)

- Model:  $f_{w,b} = wx + b$
- Parameters:  $w, b$
- Cost function:
$$\mathcal{L}(w, b) = \frac{1}{2m} \sum_{i=1}^m \left( f_{w,b}(x^{(i)}) - y^{(i)} \right)^2$$
- Minimization problem:  $\min_{w, b \in \mathbb{R}} \mathcal{L}(w, b)$


Simplified case:  $b = 0$

$$f_{w,b=0}(x) := f_w(x) = wx$$


$$\mathcal{L}(w) = \frac{1}{2m} \sum_{i=1}^m \left( f_w(x^{(i)}) - y^{(i)} \right)^2$$

$$\longrightarrow \min_w \mathcal{L}(w)$$

## Cost function: Intuition (again!)




$$\begin{aligned}\mathcal{L}(w) &= \frac{1}{2m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2 \\ &= \frac{1}{2m} \sum_{i=1}^m (wx^{(i)} - y^{(i)})^2 \\ &= \frac{1}{2 \times 3} (0^2 + 0^2 + 0^2) = 0\end{aligned}$$




$$\mathcal{L}(w = 1, b = 0) = 0$$

## Cost function: Intuition (again!)




$$\begin{aligned}\mathcal{L}(w) &= \frac{1}{2m} (wx^{(i)} - y^{(i)})^2 \\ &= \frac{1}{2 \times 3} [(0.5 - 1)^2 + (1 - 2)^2 + (1.5 - 3)^2] \\ &= 7/12\end{aligned}$$




$$\mathcal{L}(w = 0.5, b = 0) = 0.58333$$

## Cost function: Intuition (again!)




$$\begin{aligned}\mathcal{L}(w) &= \frac{1}{2m} (wx^{(i)} - y^{(i)})^2 \\ &= \frac{1}{2 \times 3} [(0 - 1)^2 + (0 - 2)^2 + (0 - 3)^2] \\ &= 14/6\end{aligned}$$




$$\mathcal{L}(w = 0, b = 0) = 14/3 \approx 2.3333$$

# Cost function: Intuition (again!)



$$\begin{aligned}
 \mathcal{L}(w) &= \frac{1}{2 \times 3} [(w-1)^2 + (2w-2)^2 + (3w-3)^2] \\
 &= \frac{1}{6} (1+4+9)(w-1)^2 \\
 &= \frac{7}{3} (w-1)^2
 \end{aligned}$$



$$J(b=0) = \frac{7}{3} (w-1)^2$$

## Visualize cost function

For our specific example with three data points  $\{\mathbf{p}^{(1)} = (1, 1), \mathbf{p}^{(2)} = (2, 2), \mathbf{p}^{(3)} = (3, 3)\}$ :

$$\begin{aligned}\mathcal{L}(w, b) &= \frac{1}{2 \times 3} [(w \times 1 + b - 1)^2 + (w \times 2 + b - 2)^2 + (w \times 3 + b - 3)^2] \\ &= \frac{7}{3}(w - 1)^2 + \frac{b^2}{2} + 2(w - 1)b\end{aligned}$$

For general cost function using linear regression model

$$\begin{aligned}\mathcal{L}(w, b) &= \frac{1}{2m} \sum_{i=1}^m \left( \underbrace{w x^{(i)}}_{\in \mathbb{R}} + \underbrace{b}_{\in \mathbb{R}} - \underbrace{y^{(i)}}_{\in \mathbb{R}} \right)^2 \\ &= \sum_{i=1}^m (A_i w^2 + B_i b^2 + C_i w b + D_i), \quad \text{where } A_i, B_i, C_i, D_i \text{ are all constants}\end{aligned}$$

*No matter how complicated the data would be,  $\mathcal{L}(w, b)$  is a quadratic function of  $(w, b)$ :*


$$\mathcal{L}(w, b) = \gamma_{11}w^2 + \gamma_{12}wb + \gamma_{21}bw + \gamma_{22}b^2 + \gamma_{00},$$

$$\gamma_{11}, \gamma_{12}, \gamma_{21}, \gamma_{22}, \gamma_{00} \in \mathbb{R}$$




## Visualize cost function

$$f(x_1, x_2) = x_1^2 + x_2^2$$




(a) Bowl shape quadratic function  $x_1^2 + x_2^2$




(b) Contour plot of quadratic function  $x_1^2 + x_2^2$

## Visualize cost function

$$\mathcal{L}(w, b) = \frac{7}{3}(w - 1)^2 + \frac{b^2}{2} + 2(w - 1)b$$



(a) Bowl shape cost function  $\mathcal{L}(w, b)$



(b) Contour plot of  $\mathcal{L}(w, b)$

## Find minimization of a function


- Have some function  $f(x_1, x_2, \dots, x_n)$
- Want to minimize function  $f$  w.r.t.  $(x_1, x_2, \dots, x_n)$ :  $\min_{x_1, \dots, x_n} f(x_1, \dots, x_n)$
- The stationary point  $\mathbf{x} = (x_1^*, \dots, x_n^*)$  is the solution of

$$\nabla f(x_1, \dots, x_n) = \mathbf{0} \quad \Longleftrightarrow \quad \begin{cases} \frac{\partial f}{\partial x_1} = 0 \\ \dots \\ \frac{\partial f}{\partial x_n} = 0 \end{cases}$$


Last example:  $\mathcal{L}(w, b) = \frac{7}{3}(w-1)^2 + \frac{b^2}{2} + 2(w-1)b$

$$\begin{cases} 0 = \frac{\partial J}{\partial w} = \frac{2}{3}(-7 + 3b + 7w) \\ 0 = \frac{\partial J}{\partial b} = -2 + b + 2w \end{cases} \Rightarrow \begin{cases} w = 1 \\ b = 0 \end{cases}$$

## Find minimization of a function: Stationary point, 1D example




$$f(x) = x^3 + 6x^2 + 8x$$




$$f'(x) = 0 \rightarrow x = x_{\max}, x = x_{\min}$$

## Find minimization of a function: Stationary point, 2D example

$$x_1^2 + x_2^2$$



$$x_1^2 + x_2^2$$



# Gradient descent

## Problem setting

- Have some function  $f(\mathbf{x}) = f(x_1, \dots, x_n)$
- Want to minimize function  $f$  w.r.t.  
 $\mathbf{x} = (x_1, \dots, x_n)$

## Simple idea

- Start with some initialization  $(x_1^{[0]}, \dots, x_n^{[0]})$
- Update  $\mathbf{x} = (x_1, \dots, x_n)$  to reduce  $f(x_1, \dots, x_n)$
- Until we are already at or near a *minimum*





Figure: Gradient descent

## Gradient descent: key idea – how to change function value at fastest rate?


Stand at  $(x_1^{[*]}, x_2^{[*]})$ , what direction  $\mathbf{n}$  ( $|\mathbf{n}| = 1$ ) would give the fastest change in the value  $f$ ?



Contour plot of  $f(x_1, x_2) = x_1^2 + x_2^2$

# Gradient descent: key idea – how to change function value at fastest rate?

Stand at  $(x_1^{[*]}, x_2^{[*]})$ , what direction  $\mathbf{n}$  ( $|\mathbf{n}| = 1$ ) would give the fastest change in the value  $f$ ?




Contour plot of  $f(x_1, x_2) = x_1^2 + x_2^2$



# Gradient descent: key idea – how to change function value at fastest rate?

Stand at  $(x_1^{[*]}, x_2^{[*]})$ , what direction  $\mathbf{n}$  ( $|\mathbf{n}| = 1$ ) would give the fastest change in the value  $f$ ?




Contour plot of  $f(x_1, x_2) = x_1^2 + x_2^2$

- $f$  at  $(x_1^{[*]}, x_2^{[*]})$  changes fastest in the direction, called here  $\mathbf{n}$  ( $|\mathbf{n}| = 1$ ), orthogonal to the contour line  $\mathcal{C}$  going through  $(x_1^{[*]}, x_2^{[*]})$
- The gradient of  $f$  at  $(x_1^{[*]}, x_2^{[*]})$  is also orthogonal to the contour line  $\mathcal{C}$ .<sup>a</sup>

# Gradient descent: key idea – how to change function value at fastest rate?

Stand at  $(x_1^{[*]}, x_2^{[*]})$ , what direction  $\mathbf{n}$  ( $|\mathbf{n}| = 1$ ) would give the fastest change in the value  $f$ ?



Contour plot of  $f(x_1, x_2) = x_1^2 + x_2^2$

➤  $f$  at  $(x_1^{[*]}, x_2^{[*]})$  changes fastest in the direction, called here  $\mathbf{n}$  ( $|\mathbf{n}| = 1$ ), orthogonal to the contour line  $\mathcal{C}$  going through  $(x_1^{[*]}, x_2^{[*]})$

➤ The gradient of  $f$  at  $(x_1^{[*]}, x_2^{[*]})$  is also orthogonal to the contour line  $\mathcal{C}$ .<sup>a</sup>

➡  $f$  changes fastest in the direction of the gradient of  $f$ :

$$\mathbf{n} = \frac{\nabla f(x_1^{[*]}, x_2^{[*]})}{|\nabla f(x_1^{[*]}, x_2^{[*]})|} = \frac{1}{|\nabla f|} \left( \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right)$$

<sup>a</sup>Please look at the supplementary note if you are interested.

# Gradient descent algorithm

Repeat until convergence

$$\begin{cases} x_1^{[t+1]} = x_1^{[t]} - \alpha^{[t]} \frac{\partial}{\partial x_1} f(x_1^{[t]}, \dots, x_n^{[t]}) \\ \dots \\ x_n^{[t+1]} = x_n^{[t]} - \alpha^{[t]} \frac{\partial}{\partial x_n} f(x_1^{[t]}, \dots, x_n^{[t]}) \end{cases} \Leftrightarrow \mathbf{x}^{[t]} = \mathbf{x}^{[t]} - \alpha^{[t]} \nabla f(\mathbf{x}^{[t]}), \quad t = 0, \dots, \infty \quad (3)$$

$\frac{\partial}{\partial x_1} f(x_1, \dots, x_n)$	(partial) derivatives
$\alpha^{[t]}$	learning rate

- We must simultaneously update  $\mathbf{x} = (x_1, \dots, x_n)$ .
- We can keep or change  $\alpha^{[t]}$  at iteration step.

# Gradient descent algorithm

**Correct:** Simultaneous update

$$\text{tmp}x_1 = x_1 - \alpha \frac{\partial f}{\partial x_1}(x_1, x_2)$$

$$\text{tmp}x_2 = x_2 - \alpha \frac{\partial f}{\partial x_2}(x_1, x_2)$$

$$x_1 = \text{tmp}x_1$$

$$x_2 = \text{tmp}x_2$$

□ Python code: *Correct*

```
df_dx1 = compute_df_dx1(x1, x2)
df_dx2 = compute_df_dx2(x1, x2)
tmpx1 = x1 - df_dx1
tmpx2 = x2 - df_dx2
x1 = tmpx1
x2 = tmpx2
```

**Incorrect:** Update before compute deri.

$$\text{tmp}x_1 = x_1 - \alpha \frac{\partial}{x_1}$$

$$x_1 = \text{tmp}x_1$$

$$\text{tmp}x_2 = x_2 - \alpha \frac{\partial f}{\partial x_2}(\underbrace{x_1}_{\text{tmp}x_1}, x_2)$$

□ Python code: *Incorrect*

```
df_dx1 = compute_df_dx1(x1, x2)
tmpx1 = x1 - df_dx1
x1 = tmpx1  # don't do this!
# don't do this, x1 holds values tmpx1 now
df_dx2 = compute_df_dx2(x1, x2)
tmpx2 = x2 - df_dx2
x2 = tmpx2
```

# Gradient descent algorithm

Coming back to our cost function  $J = \mathcal{L}(w, b)$ :  $f \rightarrow J$ ,  $x_1 \rightarrow w$ ,  $x_2 \rightarrow b$


Repeat until convergence:

$$\begin{cases} w = w - \alpha \frac{\partial J}{\partial w}(w, b) \\ b = b - \alpha \frac{\partial J}{\partial b}(w, b) \end{cases} \quad (4)$$

❑ **Python code:** This is not efficient, just for illustration

```
dJ_dw = compute_dJ_dw(w, b)
dJ_db = compute_dJ_db(w, b)
tmp_w = w - dJ_dw
tmp_b = w - dJ_db
w = tmp_w
b = tmp_b
```

## Gradient descent: why minus sign in front of $\alpha$




If  $\frac{df}{dx}(x_0) > 0$ ,  $f$  is ↗ in a neighbor  $(x_0 - \delta, x_0 + \delta)$ .

To go in the opposite direction of  $f$  ↗,  $x_0$  ↘

$$x_{\text{next}} = x_0 - \underbrace{\alpha}_{>0} \cdot \underbrace{df/dx}_{>0} < x_0$$


# Gradient descent: why minus sign in front of $\alpha$



If  $\frac{df}{dx}(x_0) > 0$ ,  $f$  is ↗ in a neighbor  $(x_0 - \delta, x_0 + \delta)$ .

To go in the opposite direction of  $f$  ↗,  $x_0$  ↘

$$x_{\text{next}} = x_0 - \underbrace{\alpha}_{>0} \cdot \underbrace{df/dx}_{>0} < x_0$$



If  $\frac{df}{dx}(x_0) < 0$ ,  $f$  is ↘ in a neighbor  $(x_0 - \delta, x_0 + \delta)$ .

To go in the opposite direction of  $f$  ↘,  $x_0$  ↗

$$x_{\text{next}} = x_0 - \underbrace{\alpha}_{>0} \cdot \underbrace{df/dx}_{<0} > x_0$$

## Gradient descent: extra thought

How about we use the plus sign in front of  $\alpha$ ?

Where would it lead to?




## Gradient descent: extra thought

That's true!

It is the algorithm for solving the maximization problem.

## Learning rate $\alpha$

Let us consider  $x = x - \alpha \frac{df}{dx}$  with  $f(x) = (x - 1)^2$  and  $x^{[0]} = 0.1$ .



$$\alpha = 0.025, x^{[14]} = 0.5830$$


$$\alpha = 0.05, x^{[14]} = 0.8147$$

$$\alpha = 0.1, x^{[14]} = 0.9683$$


If  $\alpha$  is too small, gradient descent may be slow.

## Learning rate $\alpha$


Let us consider  $x = x - \alpha \frac{df}{dx}$  with  $f(x) = (x - 1)^2$  and  $x^{[0]} = 0.1$ . We now increase  $\alpha$  as follows



$$\alpha = 0.5, x^{[24]} = 0.5830$$




$$\alpha = 0.9, x^{[24]} = 0.8147$$



$$\alpha = 1.0, x^{[24]} = 0.9683$$

## Learning rate $\alpha$

➡ Let us initialize  $x^{[0]} = 0.95$ , much close to the solution  $x_{\min} = 1$  but then use  $\alpha = 1.1$



➡ If  $\alpha$  is too large, gradient descent may

- overshoot, never reach minimum
- fail to converge, diverge
- or be too slow too (see last slide – it just shoots over back and forth)

## Batch gradient descent

**Batch:** Each step of gradient descent uses **all the training examples**.

**Batch** gradient descent uses **all the training examples** so that we calculate the derivatives of  $\mathcal{L}(w, b)$  w.r.t.  $w$  and  $b$  by summing up over all training examples.

$$\partial J / \partial w = \sum_{\text{all training samples}} \dots, \quad \partial J / \partial b = \sum_{\text{all training samples}} \dots$$

	$x$	$y$
	BMI	RI
(1)	32.1	151
(2)	21.6	75
(3)	30.5	141
(4)	25.3	206
$\vdots$	$\vdots$	$\vdots$
( $m$ )	23	135

*‘Other’ gradient descent:* Each step of gradient descent uses just subsets of the training examples. So, we use “approximate” partial derivatives of  $\mathcal{L}(w, b)$  by summing up over a portion of training examples. We call this **minibatch gradient descent**.

$$\partial J / \partial w = \sum_{\text{subset of samples}} \dots, \quad \partial J / \partial b = \sum_{\text{subset of samples}} \dots$$

# Training set versus test set

- **Training set:**

Set of examples used for fitting/training the regression models

- **Test set:**

Set of examples used for assessing how well the regression models perform or generalize to the unseen data (test data)

➡ We will learn in the next lectures some metrics to evaluate how well a regression model performs on a given data set:

- R squared score/coefficient of determination
- Mean squared error
- Accuracy score (for classification problem)